
Redes de Computadores


Shortest Paths in Networks

Manuel P. Ricardo, Rui Prior


Universidade do Porto

-
- » *What is a graph?*
 - » *What is a spanning tree?*
 - » *What is a shortest path tree?*
 - » *How are paths defined in a network?*
 - » *How does the Dijkstra algorithm work?*
 - » *How does a link state routing protocol work?*
 - » *How does a node learn about neighbours?*
 - » *How does the Bellman-Ford algorithm work?*
 - » *How does a distance vector work?*
 - » *What are the limitations of the layer 2 network of switches?*
 - » *How does the IEEE spanning tree protocol work?*
 - » *What is the maximum capacity of a flow network?*

Graph – Directed and Undirected



a) Directed graph



b) Undirected graph

$$G = (V, E)$$

$$V = \{v_1, v_2, v_3, v_4, v_5, v_6\}, \quad |V| = 6$$

$$E = \{(v_1, v_2), (v_2, v_3), (v_2, v_4), (v_3, v_4), \\ (v_4, v_3), (v_3, v_5), (v_4, v_5), (v_5, v_6)\}, \quad |E| = 8$$


$$G = (V, E)$$

$$V = \{v_1, v_2, v_3, v_4, v_5, v_6\}, \quad |V| = 6$$

$$E = \{(v_1, v_2), (v_2, v_3), (v_2, v_4), (v_3, v_4), \\ (v_3, v_5), (v_4, v_5), (v_5, v_6)\}, \quad |E| = 7$$

Tree


- ♦ Tree $T = (V, E)$
 - » connected graph with no cycles
 - » number of edges $|E| = |V| - 1$
 - » any two vertices of the tree are connected by exactly one path
- ♦ A tree T is said to span a graph $G = (V, E)$ (spanning tree) if
 - » $T = (V, E')$ and $E' \subseteq E$



Shortest Path Trees

- ♦ Graphs and Trees can be weighted

- » $G=(V, E, w)$
- » $T=(V, E', w)$
- » $w: E \rightarrow \mathbb{R}$




- ♦ Total cost of a tree $T \rightarrow C_{total}(T) = \sum_{i=1}^{|E|} w(e_i)$
- ♦ Minimum Spanning Tree $T^* \rightarrow C_{total}(T^*) = \min(C_{total}(T))$
 - » algorithms used to compute MST: Prim, Kruskal
- ♦ **Shortest Path Tree (SPT) Rooted at Vertex s**
 - » tree composed by the **union of the shortest paths between s and each of other vertices of G**
 - » algorithms used to compute SPT: **Dijkstra, Bellman-Ford**
 - » depends on s
- ♦ Computer networks use **Shortest Path Trees**

Routing in Layer 3 Networks

Forwarding, Routing

- ♦ Forwarding → data plane
 - » directing packet from input to output link
 - » using a forwarding table
- ♦ Routing → control plane
 - » computing paths the packets will follow
 - » routers exchange messages
 - » each router creates its forwarding table



Importance of Routing

- ◆ End-to-end performance
 - » path affects quality of service
 - » delay, throughput, packet loss


- ◆ Use of network resources
 - » balance traffic over routers and links
 - » avoiding congestion by directing traffic to less-loaded links

- ◆ Transient disruptions
 - » failures, maintenance
 - » limiting packet loss and delay during changes

Shortest-Path Routing


Path-selection model

- » Destination-based
- » Load-insensitive (e.g., static link weights)
- » Minimum sum of link weights
 - Special case: minimum hop count (all weights equal)



Shortest-Path Problem

- ◆ Given a network topology with link costs
 - » $c(x,y)$ - link cost from node x to node y
 - » infinity if x and y are not direct neighbors
- ◆ Compute the least-cost paths from source u to all nodes
 - $p(v)$ - node predecessor of node v in the path to u



Dijkstra's Shortest-Path Algorithm

- ♦ Iterative algorithm
 - » After k iterations \rightarrow known least-cost paths to k nodes
- ♦ $\mathbf{S} \rightarrow$ set of nodes for which least-cost path is known
 - » Initially, $\mathbf{S} = \{\mathbf{u}\}$, where \mathbf{u} is the source node
 - » Add one node to \mathbf{S} in each iteration
- ♦ $\mathbf{D}(\mathbf{v}) \rightarrow$ current cost of path from source to node \mathbf{v}
 - » Initially
 - $\mathbf{D}(\mathbf{v}) = \mathbf{c}(\mathbf{u}, \mathbf{v})$ for all nodes \mathbf{v} adjacent to \mathbf{u}
 - $\mathbf{D}(\mathbf{v}) = \infty$ for all other nodes \mathbf{v}
 - » Continually update $\mathbf{D}(\mathbf{v})$ when shorter paths are learned

Dijkstra's Algorithm

1. *Initialization:*

2. $S = \{u\}$

3. for all nodes v

4. if v adjacent to u

5. $D(v) = c(u,v)$

6. $p(v) = u$

7. else $D(v) = \infty$

8.

9. *Loop:*

10. find node w not in S with the smallest $D(w)$

11. add w to S

12. update $D(v)$, $p(v)$ for all v adjacent to w and not in S :





13. if $D(w) + c(w,v) < D(v)$

14. $D(v) = D(w) + c(w,v)$





15. $p(v) = w$

16. *repeat until all nodes are in S*

Dijkstra's Algorithm - Example




Dijkstra's Algorithm - Example



Shortest-Path Tree

◆ Shortest-path tree from u

◆ Forwarding table at u



NOTE: keeping track of the predecessors is necessary to build the shortest paths tree


	link
v	(u,v)
w	(u,w)
x	(u,w)
y	(u,v)
z	(u,v)
s	(u,w)
t	(u,w)

Link-State Routing

- ◆ Each router keeps track of its incident links
 - » link up, link down
 - » cost on the link
- ◆ Each router broadcasts link states
 - » every router gets a complete view of the graph
- ◆ Each router runs Dijkstra's algorithm, to
 - » compute the shortest paths
 - » construct the forwarding table
- ◆ Example protocols
 - » Open Shortest Path First (OSPF)
 - » Intermediate System – Intermediate System (IS-IS)

Detection of Topology Changes





- ◆ Beacons generated by routers on links
 - » Periodic “hello” messages in both directions
 - » After a few missed “hellos” → link failure detected



Broadcasting the Link State


- ◆ How to Flood the link state?
 - » every node sends link-state information through adjacent links
 - » next nodes forward that info to all links except the one where the information arrived

- ◆ When to initiate flooding?
 - » Topology change
 - link or node failure/recovery
 - link cost change
 - » Periodically
 - refresh link-state information
 - typically 30 minutes




Scaling Link-State Routing

- ◆ Overhead of link-state routing
 - » flooding link-state packets throughout the network
 - » running Dijkstra's shortest-path algorithm
- ◆ Introducing hierarchy through “areas”



Bellman-Ford Algorithm

- ◆ Define distances at each node x
 - » $d_x(y) = \text{cost of least-cost path from } x \text{ to } y$
- ◆ Update distances based on neighbors
 - » $d_x(y) = \min \{c(x,v) + d_v(y)\}$ over all neighbors v



$$d_u(z) = \min \{ c(u,v) + d_v(z), \\ c(u,w) + d_w(z) \}$$

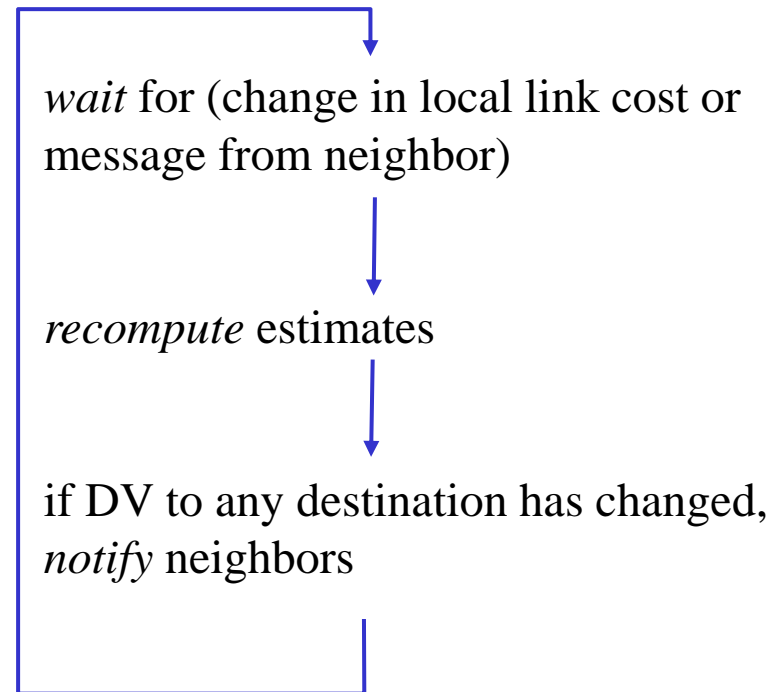
Distance Vector Algorithm

- ♦ $c(x,v)$ = cost for direct link from x to v
node x maintains costs of direct links $c(x,v)$
- ♦ $D_x(y)$ = estimate of least cost from x to y
node x maintains distance vector $\mathbf{D}_x = [D_x(y): y \in N]$
- ♦ Node x maintains also its neighbors' distance vectors
for each neighbor v , x maintains $\mathbf{D}_v = [D_v(y): y \in N]$
- ♦ Each node v periodically sends \mathbf{D}_v to its neighbors
 - » and neighbors update their own distance vectors
 - » $D_x(y) \leftarrow \min_v \{c(x,v) + D_v(y)\}$ for each node $y \in N$
- ♦ Over time, the distance vector \mathbf{D}_x converges

Distance Vector Algorithm

- ◆ Iterative, asynchronous
 - each local iteration caused by:
 - local link cost change
 - distance vector update message from neighbor
- ◆ Distributed
 - » node notifies neighbors only when its DV changes
- ◆ Neighbors then notify their neighbors, if necessary

Each node:



Distance Vector Example - Step 0

Table for A			Table for B		
Dst	Cst	Hop	Dst	Cst	Hop
A	0	(A)	A	4	A
B	4	B	B	0	(B)
C	∞	—	C	∞	—
D	∞	—	D	3	D
E	2	E	E	∞	—
F	6	F	F	1	F




Table for C			Table for D			Table for E			Table for F		
Dst	Cst	Hop	Dst	Cst	Hop	Dst	Cst	Hop	Dst	Cst	Hop
A	∞	—	A	∞	—	A	2	A	A	6	A
B	∞	—	B	3	B	B	∞	—	B	1	B
C	0	(C)	C	1	C	C	∞	—	C	1	C
D	1	D	D	0	(D)	D	∞	—	D	∞	—
E	∞	—	E	∞	—	E	0	(E)	E	3	E
F	1	F	F	∞	—	F	3	F	F	0	(F)

Distance Vector Example - Step 1

Table for A			Table for B		
Dst	Cst	Hop	Dst	Cst	Hop
A	0	(A)	A	4	A
B	4	B	B	0	(B)
C	7	F	C	2	F
D	7	B	D	3	D
E	2	E	E	4	F
F	5	E	F	1	F




Table for C			Table for D			Table for E			Table for F		
Dst	Cst	Hop	Dst	Cst	Hop	Dst	Cst	Hop	Dst	Cst	Hop
A	7	F	A	7	B	A	2	A	A	5	B
B	2	F	B	3	B	B	4	F	B	1	B
C	0	(C)	C	1	C	C	4	F	C	1	C
D	1	D	D	0	(D)	D	∞	—	D	2	C
E	4	F	E	∞	—	E	0	(E)	E	3	E
F	1	F	F	2	C	F	3	F	F	0	(F)

Distance Vector Example - Step 2

Table for A			Table for B		
Dst	Cst	Hop	Dst	Cst	Hop
A	0	(A)	A	4	A
B	4	B	B	0	(B)
C	6	E	C	2	F
D	7	B	D	3	D
E	2	E	E	4	F
F	5	E	F	1	F





Table for C			Table for D			Table for E			Table for F		
Dst	Cst	Hop	Dst	Cst	Hop	Dst	Cst	Hop	Dst	Cst	Hop
A	6	F	A	7	B	A	2	A	A	5	B
B	2	F	B	3	B	B	4	F	B	1	B
C	0	(C)	C	1	C	C	4	F	C	1	C
D	1	D	D	0	(D)	D	5	F	D	2	C
E	4	F	E	5	C	E	0	(E)	E	3	E
F	1	F	F	2	C	F	3	F	F	0	(F)

Routing Information Protocol (RIP)

- ◆ Distance vector protocol
 - » nodes send distance vectors every 30 seconds
 - » or, when an update causes a change in routing
- ◆ RIP is limited to small networks

BGP – The Exterior Gateway Routing Protocol



(a)

Information F receives
from its neighbors about D

From B: "I use BCD"
From G: "I use GCD"
From I: "I use IFGCD"
From E: "I use EFGCD"

(b)

(a) A set of BGP routers. (b) Information sent to F

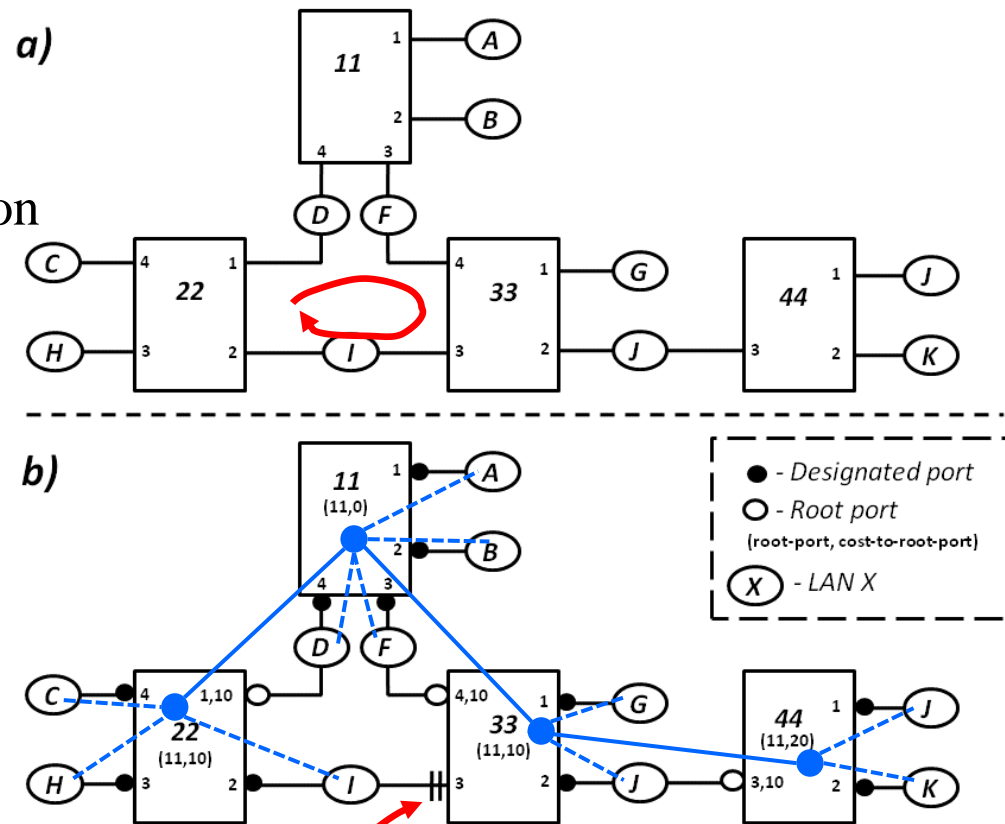
Unique Spanning Tree in Ethernet Networks

L2 Networking - Single Tree Required

- Ethernet frame
 - No hop-count
 - Could loop forever
 - Broadcast frame, misconfiguration

- Layer 2 network
 - **Required to have tree topology**
 - Single path between every pair of stations


- Spanning Tree Protocol (STP)
 - Running in bridges
 - Helps building the spanning tree
 - Blocks ports



Constructing a Spanning Tree

Distributed algorithm

- » switches need to elect a “root”
 - the switch with the smallest identifier
- » each switch identifies if its interface is on **the shortest path from the root**
- » messages (Y, d, X)
 - from node X
 - claiming Y is the root
 - and the distance is d




Steps in Spanning Tree Algorithm

- ◆ Initially, each switch thinks it is the root
 - » switch sends a message out every interface
 - » identifying itself as the root with distance 0
 - » example: switch X announces (X, 0, X)
- ◆ Other switches update their view of the root
 - » upon receiving a message, check the root id
 - » if the new id is smaller, start viewing that switch as the root
- ◆ Switches compute their distance from the root
 - » add interface *cost* to the distance received from a neighbor
 - » identify interfaces not on a shortest path to the root and exclude them from the spanning tree

Example - Switch #4's Viewpoint


- ◆ Switch #4 thinks it is the root
 - » sends (4, 0, 4) message to 2 and 7
- ◆ Then, switch #4 hears from #2
 - » receives (2, 0, 2) message from 2
 - » ... and thinks that #2 is the root
 - » and realizes it is just one hop away
- ◆ Then, switch #4 hears from #7
 - » receives (2, 1, 7) from 7
 - » and realizes this is a longer path
 - » so, prefers its own one-hop path
 - » and removes 4-7 link from the tree



For simplicity, consider all costs = 1

Example - Switch #4's Viewpoint

- ◆ Switch #2 hears about switch #1
 - » switch 2 hears (1, 1, 3) from 3
 - » switch 2 starts treating 1 as root
 - » and sends (1, 2, 2) to neighbors
- ◆ Switch #4 hears from switch #2
 - » switch 4 starts treating 1 as root
 - » and sends (1, 3, 4) to neighbors
- ◆ Switch #4 hears from switch #7
 - » switch 4 receives (1, 3, 7) from 7
 - » and realizes this is a longer path
 - » so, prefers its own three-hop path
 - » and removes 4-7 link from the tree



Maximum Flow of a Network


Flow Network Model

- ◆ **Flow network**

- » source s
- » sink t
- » nodes a, b and c

- ◆ Edges are labeled with **capacities**

- » (e.g. bit/s)



- ◆ Communication networks are not flow networks

- » they are queue networks
- » flow networks allow us to determine limit capacity values

Maximum Capacity of a Flow Network

- ♦ Max-flow min-cut theorem
 - » maximum amount of flow transferable through a network equals minimum value among all simple cuts of the network
- ♦ Cut \rightarrow split of the nodes V into two disjoint sets S and T
 - » $S \cup T = V$
 - » there are $2^{|V|-2}$ possible cuts
- ♦ Capacity of cut (S, T) :
$$c(S, T) = \sum_{(u,v) \mid u \in S, v \in T, (u,v) \in E} c(u, v)$$


Max-flow Min-cut - Example

$2^{|V|-2} = 8$ possible cuts


Cut	Vertices					$c(S, T)$	Feasibility
	s	a	b	c	t		
1	S	S	S	S	T	10	✓
2	S	S	S	T	T	19	✓
3	S	S	T	S	T	13	✓
4	S	S	T	T	T	17	✓
5	S	T	S	S	T	-	×
6	S	T	S	T	T	-	×
7	S	T	T	S	T	11	✓
8	S	T	T	T	T	12	✓

Maximum flow = 10


a) Cut 1




b) Cut 2



c) Cut 3



d) Cut 7



Homework

1. Review slides
2. Read from Tanenbaum
 - » Section 5.2 – Routing algorithms
 - » Section 4.8.3 – Spanning Tree Bridges