
Redes de Computadores

Application Layer

Manuel Ricardo, Rui Prior

Universidade do Porto

Based on slides provided by J.F Kurose and K.W. Ross for their book “Computer Networking: a Top-Down Approach”

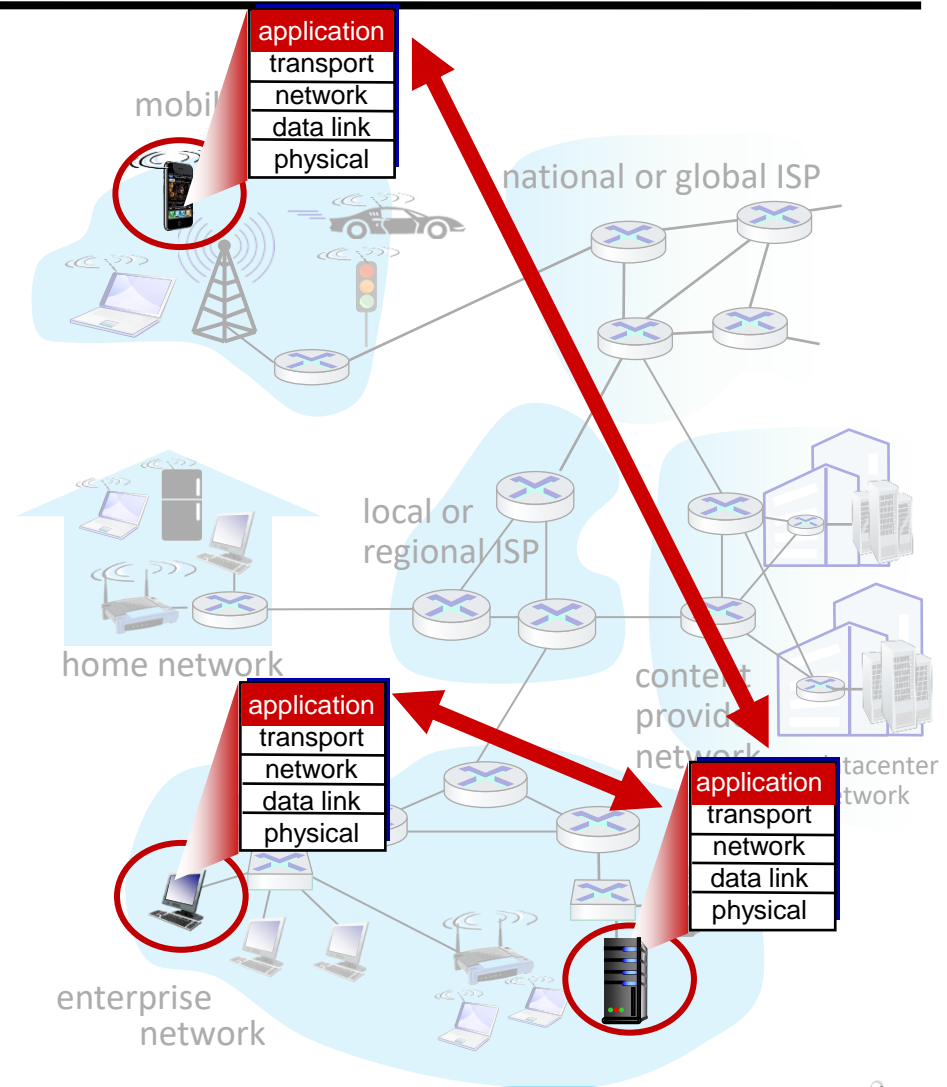
Creating a network app

write programs that

- ◆ run on (different) end systems
- ◆ communicate over network
- ◆ e.g., web server software communicates with browser software

no need to write software for network-core devices

- ◆ network-core devices do not run user applications
- ◆ applications on end systems allow for rapid app development, propagation



An application-layer protocol defines:

- ♦ **types of messages exchanged**,
 - » e.g., request, response
- ♦ **message syntax**
 - » what fields in messages & how fields are delineated
- ♦ **message semantics**
 - » meaning of information in fields
- ♦ **rules** for when and how processes send & respond to messages

open application protocols:

- ♦ defined in RFCs, everyone has access to protocol definition
- ♦ allows for interoperability
- ♦ e.g., HTTP, SMTP

proprietary application protocols:

- ♦ e.g., Skype, Zoom

A note on security

- ♦ TCP and UDP offer no security
 - » e.g., passwords are sent in cleartext, visible on the Internet
- ♦ Transport Layer Security (TLS) is an additional “layer” providing
 - » Encryption (privacy)
 - Eavesdroppers on the Internet see only “random” bits
 - » Message integrity
 - Message has not been forged or tampered with
 - » Endpoint authentication
 - Through the use of certificates
- ♦ TLS implemented as a library, included in the application layer
 - » Talks to TCP

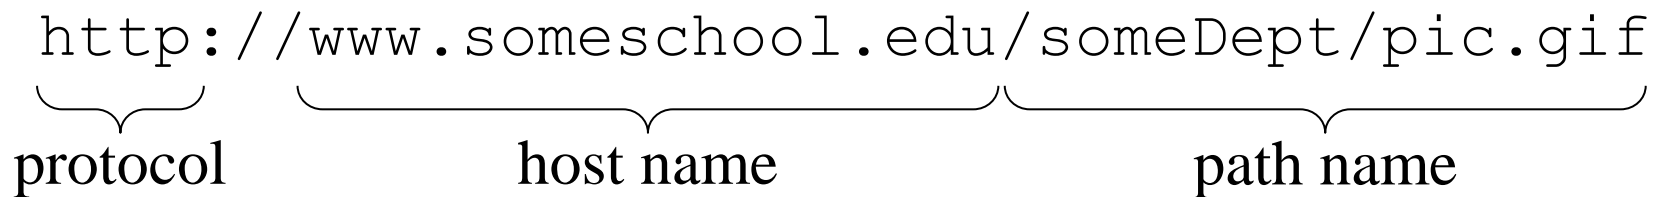
HyperText Transfer Protocol (HTTP)

Web and HTTP

First some jargon

- ♦ **Web page** consists of **objects**
- ♦ Object can be HTML file, JPEG image, Java applet, audio file,...
- ♦ Web page consists of **base HTML file** which includes several referenced objects
- ♦ Each object is addressable by a **URL**
- ♦ Example URL:

`http://www.someschool.edu/someDept/pic.gif`




protocol host name path name

HTTP overview

HTTP: hypertext transfer protocol

- ♦ Web's application layer protocol
- ♦ client/server model
 - » *client*: browser that requests, receives, “displays” Web objects
 - » *server*: Web server sends objects in response to requests



HTTP overview (continued)

Uses TCP:

- ◆ client initiates TCP connection (creates socket) to server, port 80
- ◆ server accepts TCP connection from client
- ◆ HTTP messages (application-layer protocol messages) exchanged between browser (HTTP client) and Web server (HTTP server)
- ◆ TCP connection closed

HTTP is “stateless”

- ◆ server maintains no information about past client requests

Nonpersistent HTTP

- ◆ At most one object is sent over a TCP connection

Persistent HTTP

- ◆ Multiple objects can be sent over single TCP connection between client and server

Nonpersistent HTTP

Suppose user enters URL

`www.someSchool.edu/someDepartment/home.index`

(contains text and
references to 10
jpeg images)

1a. HTTP client initiates TCP
connection to HTTP server
(process) at
`www.someSchool.edu` on port 80

1b. HTTP server at host
`www.someSchool.edu` waiting
for TCP connection at port 80.
“accepts” connection, notifying
client

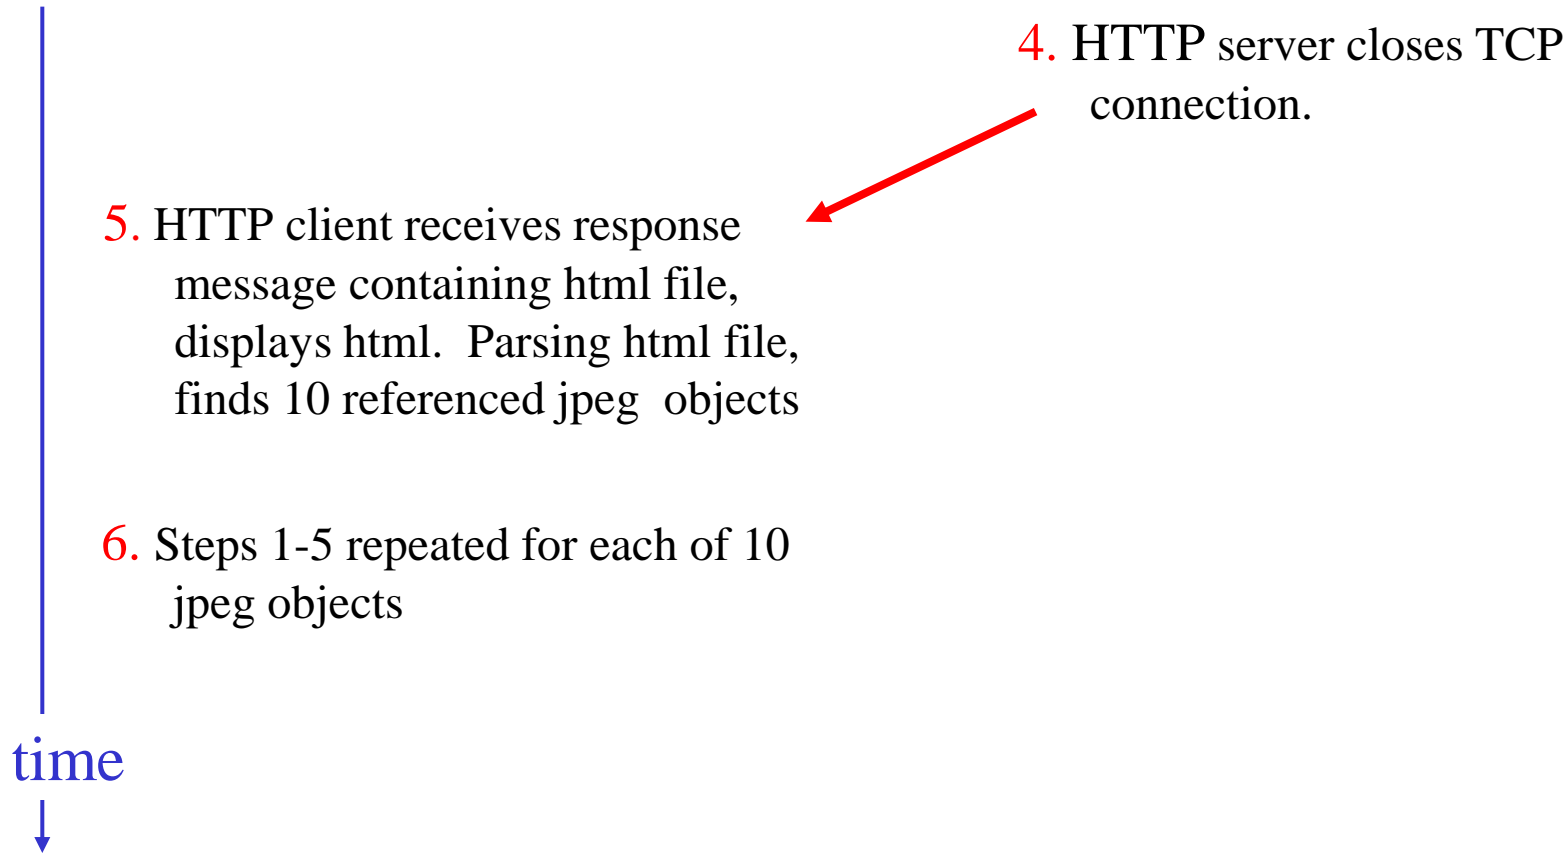
2. HTTP client sends HTTP *request
message* (containing URL) into
TCP connection socket. Message
indicates that client wants object
`someDepartment/home.index`

3. HTTP server receives request
message, forms *response message*
containing requested object, and
sends message into its socket

time



Nonpersistent HTTP (cont.)




Non-Persistent HTTP: Response time

Definition of RTT: time for a small packet to travel from client to server and back.

Response time:

- ♦ one RTT to initiate TCP connection
- ♦ one RTT for HTTP request and first few bytes of HTTP response to return
- ♦ file transmission time

total = $2RTT + \text{transmit time}$



Persistent HTTP

Nonpersistent HTTP issues:

- ◆ requires 2 RTTs per object
- ◆ OS overhead for *each* TCP connection
- ◆ browsers often open parallel TCP connections to fetch referenced objects

Persistent HTTP

- ◆ server leaves connection open after sending response
- ◆ subsequent HTTP messages between same client/server sent over open connection

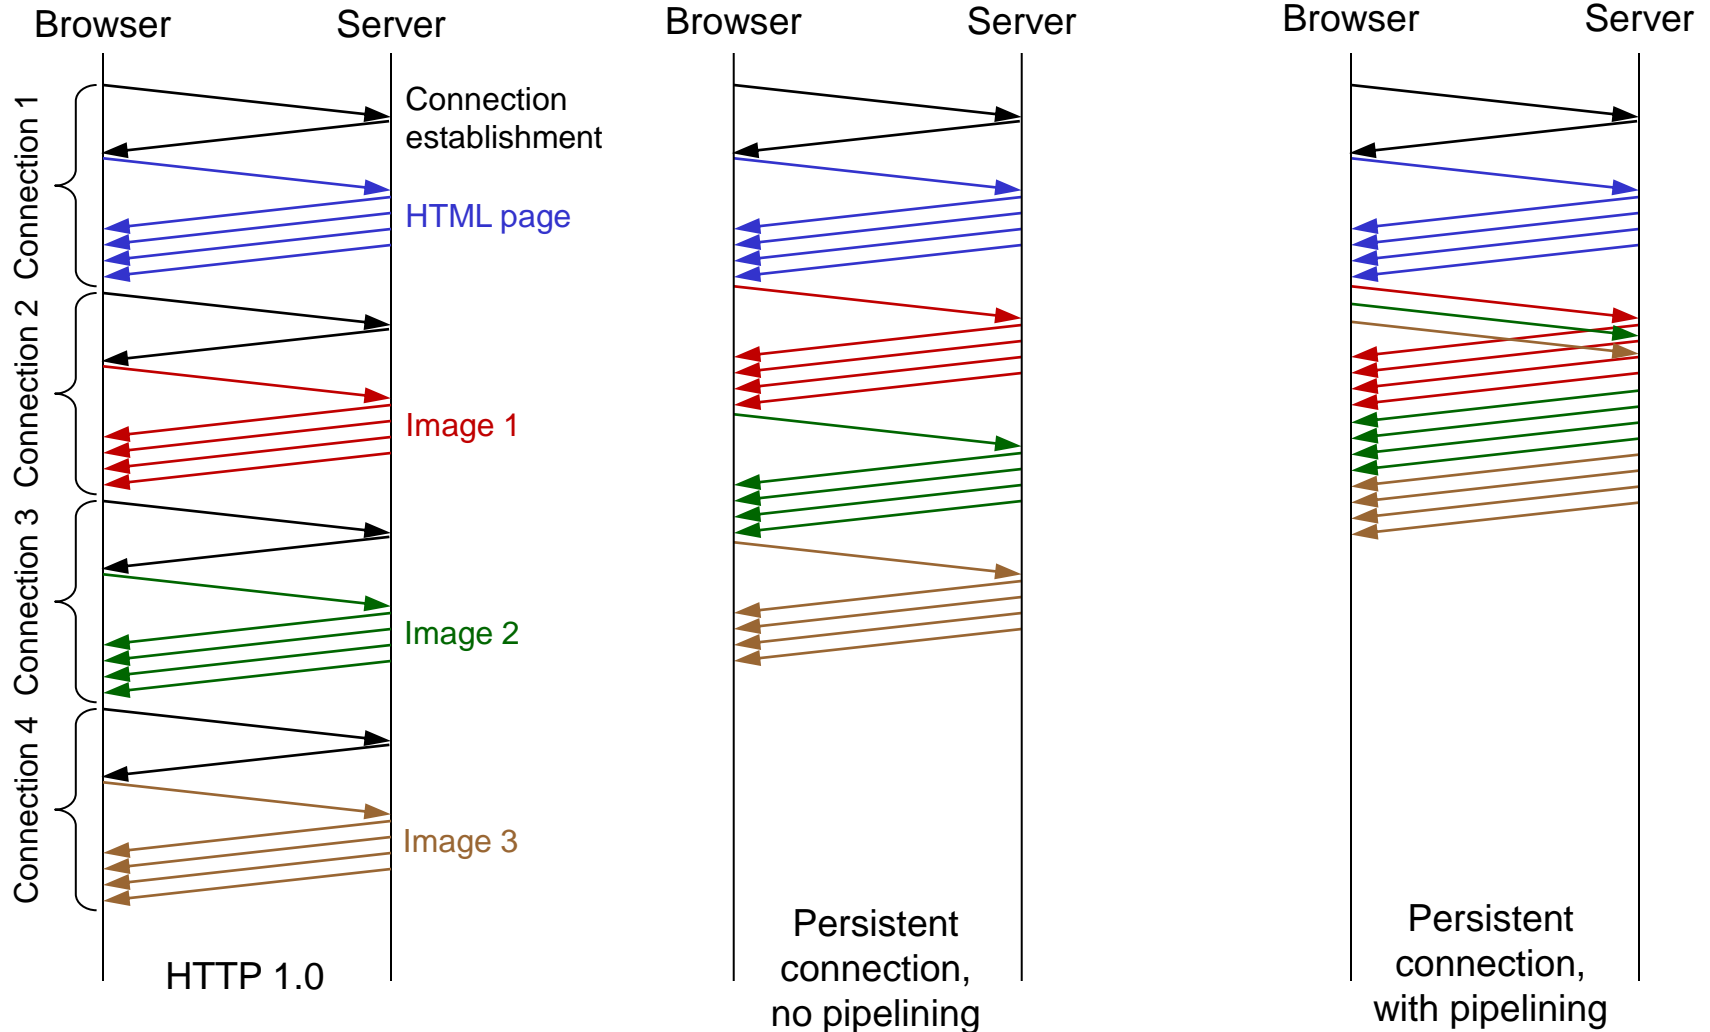
Persistent **without** pipelining:

- ◆ client makes new request only after receiving the previous reply
- ◆ one RTT per referenced object

Persistent **with** pipelining:

- ◆ default in HTTP/1.1
- ◆ client requests new object as soon as it finds its reference
- ◆ one RTT for all referenced objects

Persistent connections and pipelining



HTTP request message

- ♦ Two types of HTTP messages: *request*, *response*
- ♦ **HTTP request message:**
 - » ASCII (human-readable format)

request line
(GET, POST,
HEAD commands)

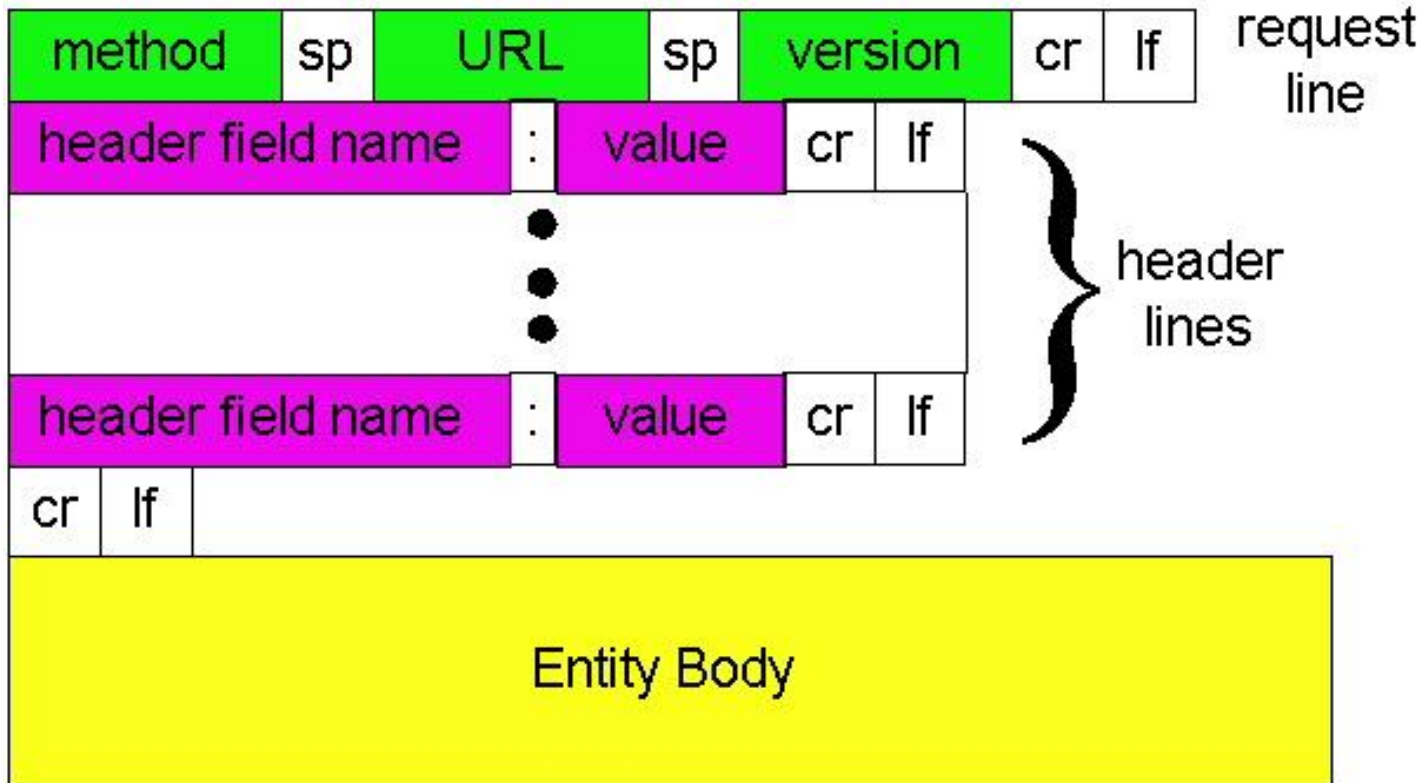
header
lines

Carriage return,
line feed
indicates end
of message

```
GET /somedir/page.html HTTP/1.1
Host: www.someschool.edu
User-agent: Mozilla/4.0
Connection: close
Accept-language: pt
```

(extra carriage return, line feed)

HTTP request message: general format



Uploading form input

POST method:

- ◆ Web page often includes form input
- ◆ Using POST method, input is uploaded to server in entity body

GET method:

- ◆ Input is uploaded in URL field of request line:

`http://www.animals.com/search?animal=monkeys&likes=banana`

Some methods

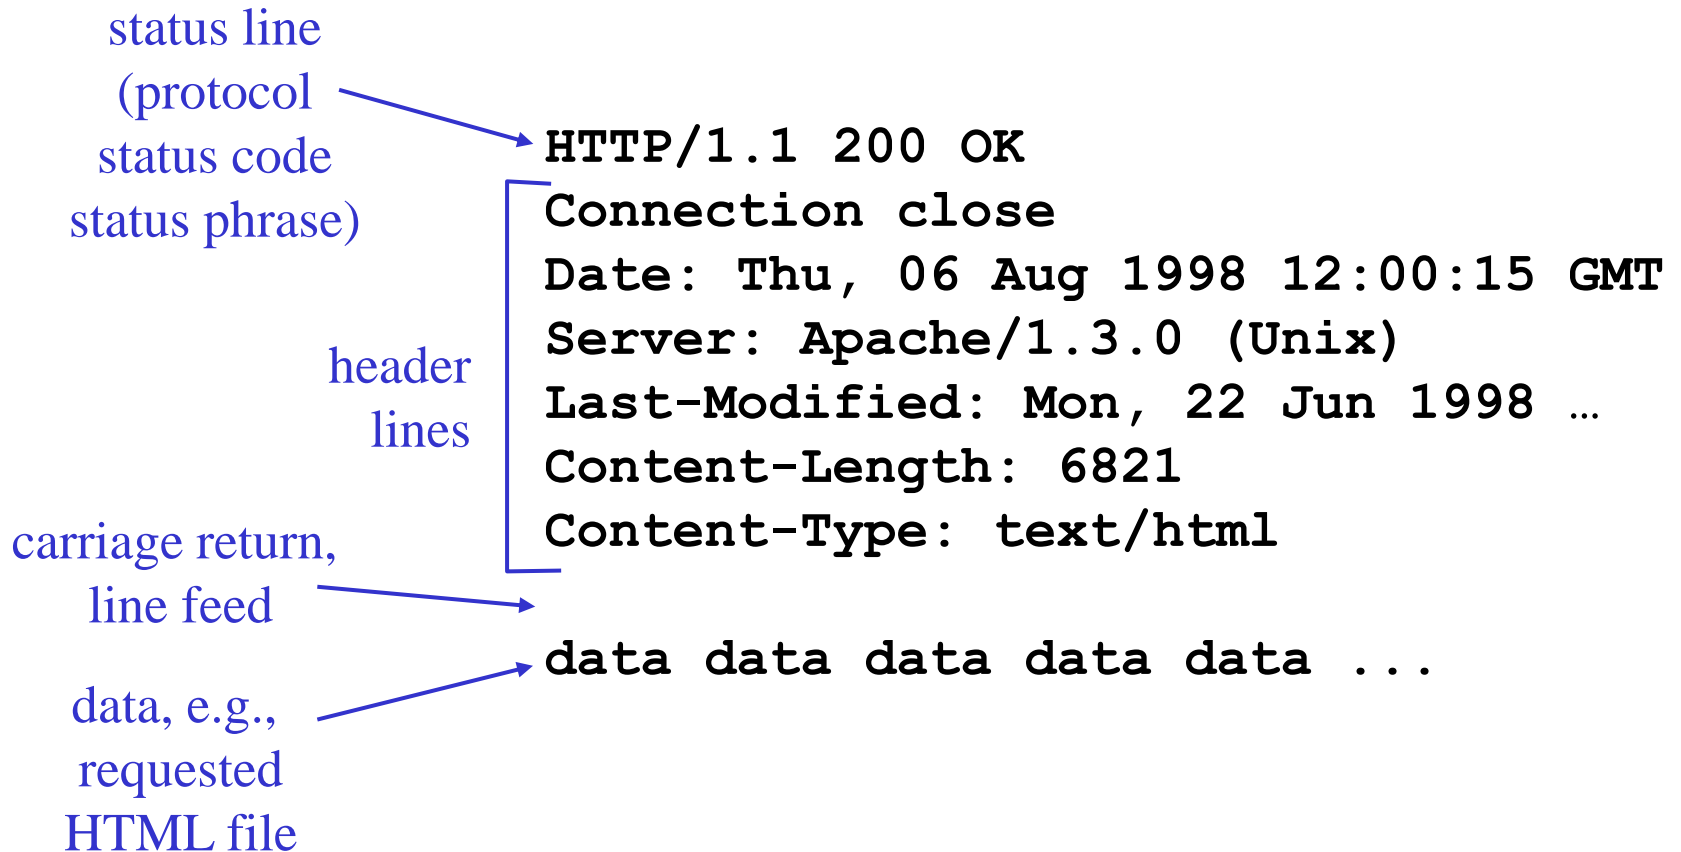
HTTP/1.0

- ♦ GET
- ♦ POST

HTTP/1.1

- ♦ GET, POST
- ♦ PUT
 - » uploads file in entity body to path specified in URL field
- ♦ DELETE
 - » deletes file specified in the URL field

HTTP response message



HTTP response status codes

In first line in server → client response message.

A few sample codes:

200 OK

- » request succeeded, requested object later in this message

301 Moved Permanently

- » requested object moved, new location specified later in this message
(Location:)

400 Bad Request

- » request message not understood by server

404 Not Found

- » requested document not found on this server

505 HTTP Version Not Supported

User-server state: cookies

- ♦ The HTTP protocol itself is stateless
- ♦ State information can be conveyed using cookies

Four components:

- 1) cookie header line of HTTP *response* message
- 2) cookie header line in HTTP *request* message
- 3) cookie file kept on user's host, managed by user's browser
- 4) back-end database at Web site

Example:

- ♦ Susan accesses the Internet from her PC
- ♦ visits specific e-commerce site for the first time
- ♦ when initial HTTP request arrives at site, site creates:
 - » unique ID
 - » entry in backend database for ID

Cookies: keeping “state”

client

server



cookie file



one week later:



usual http request msg

Amazon server
creates ID
1678 for user

usual http response
Set-cookie: 1678

usual http request msg
Cookie: 1678

cookie-
specific
action

usual http response msg

usual http request msg
Cookie: 1678


cookie-
specific
action

usual http response msg

create
entry

access

access



backend
database

Cookies (continued)

What cookies can bring:

- ◆ authorization
- ◆ shopping carts
- ◆ recommendations
- ◆ user session state (Web e-mail)

How to keep “state”:

- ◆ protocol endpoints: maintain state at sender/receiver over multiple transactions
- ◆ cookies: http messages carry state

aside

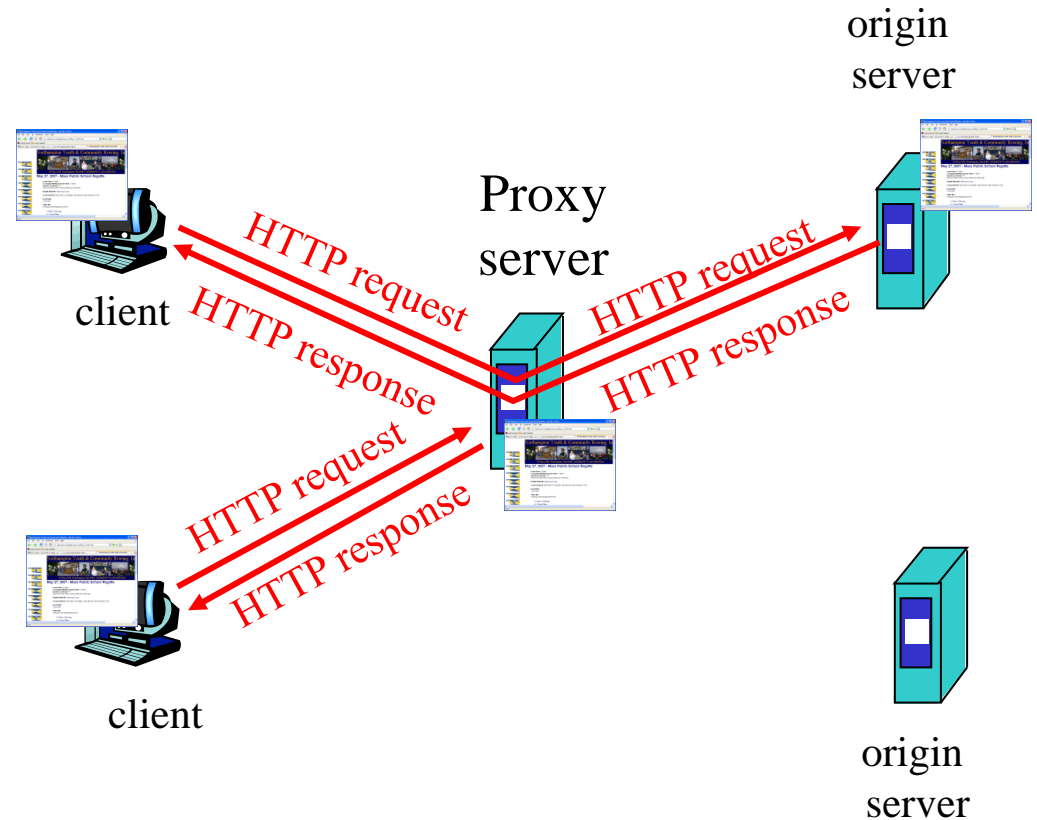
Cookies and privacy:

- ◆ cookies permit sites to learn a lot about you
- ◆ you may supply name and e-mail to sites

Web caches (proxy server)

Goal: satisfy client request without involving origin server

- ♦ user sets browser: Web accesses via cache
- ♦ browser sends all HTTP requests to cache
 - » object in cache: cache returns object
 - » else cache requests object from origin server, then returns object to client



More about Web caching

- ◆ cache acts as both client and server
- ◆ typically, cache is installed by ISP (university, company, residential ISP)

Why Web caching?

- ◆ reduce response time for client request
- ◆ reduce traffic on an institution's access link

Final notes on HTTP

Newer versions of the protocol:

- ♦ HTTP/2, HTTP/3
- ♦ Use the same messages but faster / more efficient transport

New uses of Web and HTTP:

- ♦ The Internet of Things
- ♦ REST APIs
 - » The core IETF working group
 - » RFC 6690

*Electronic Mail (e-mail):
SMTP, POP3, IMAP*

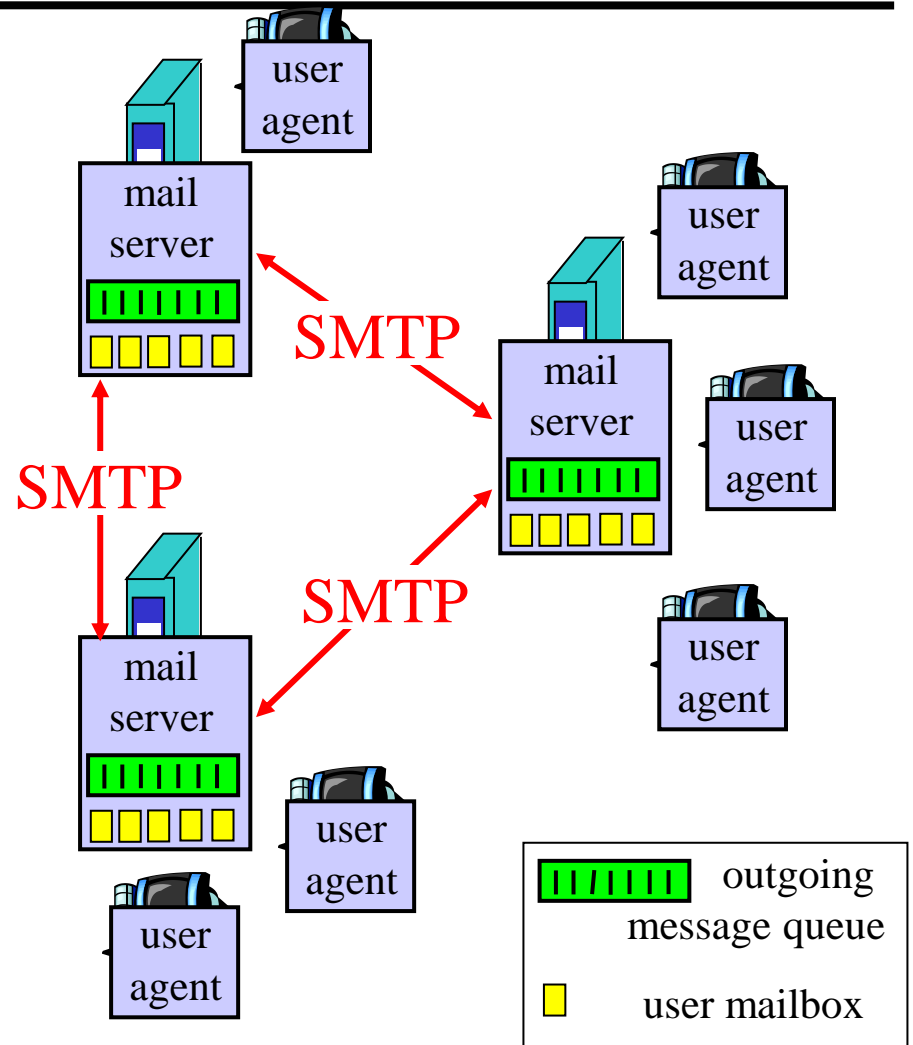
Electronic Mail

Three major components:

- ♦ user agents
- ♦ mail servers
- ♦ simple mail transfer protocol: SMTP

User Agent

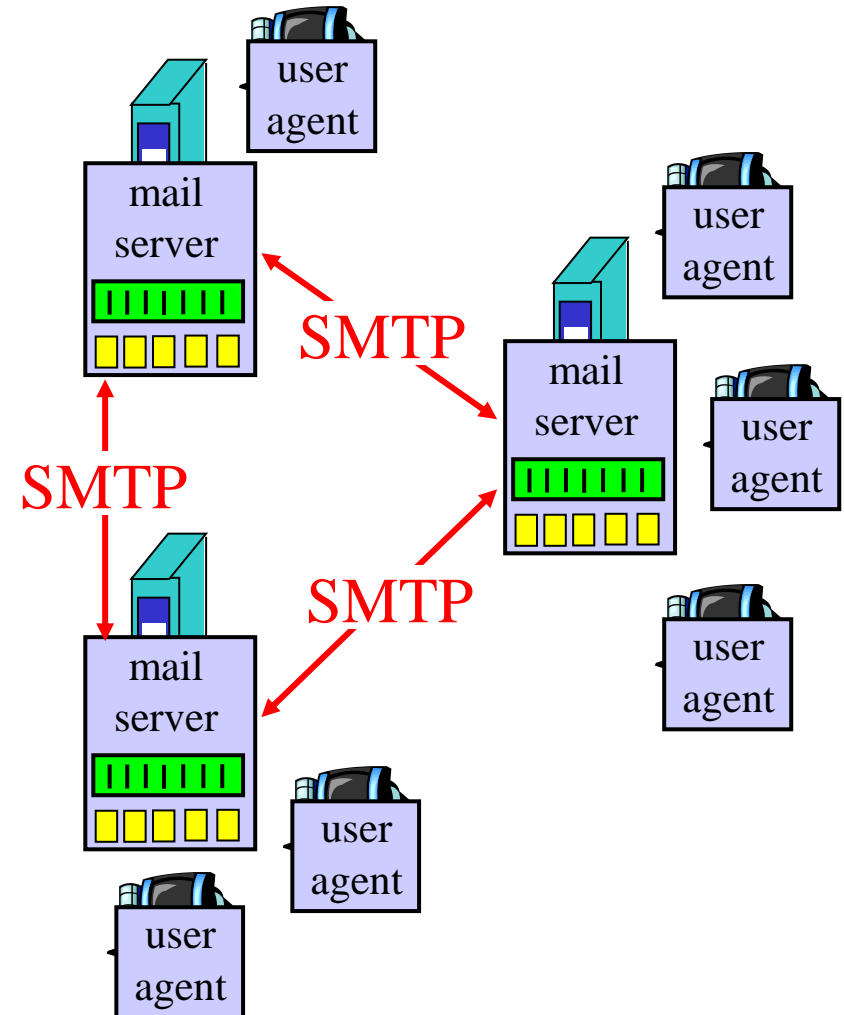
- ♦ a.k.a. “mail reader”
- ♦ composing, editing, reading mail messages
- ♦ e.g., Eudora, Outlook, elm, Mozilla Thunderbird
- ♦ outgoing, incoming messages stored on server



Electronic Mail: mail servers

Mail Servers

- ♦ **mailbox** contains incoming messages for user
- ♦ **message queue** of outgoing (to be sent) mail messages
- ♦ **SMTP protocol** between mail servers to send email messages
 - » client: sending mail server
 - » “server”: receiving mail server

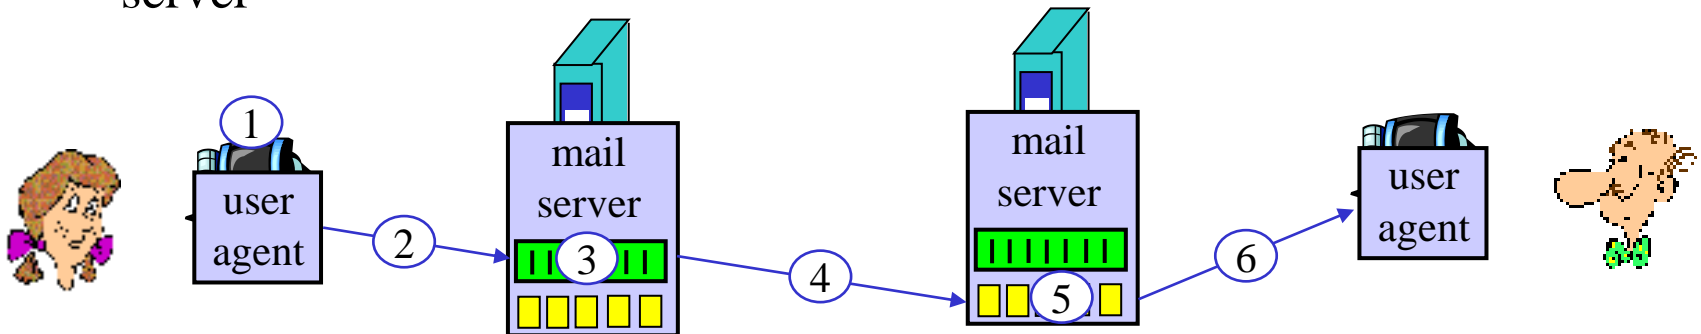


Electronic Mail: SMTP [RFC 2821]

- ♦ uses TCP to reliably transfer email message from client to server, port 25
- ♦ direct transfer: sending server to receiving server
- ♦ three phases of transfer
 - » handshaking (greeting)
 - » transfer of messages
 - » closure
- ♦ command/response interaction
 - » commands: ASCII text
 - » response: status code and phrase
- ♦ messages must be in 7-bit ASCII

Scenario: Alice sends message to Bob

- 1) Alice uses UA to compose message and “to” bob@some school . edu
- 2) Alice’s UA sends message to her mail server; message placed in message queue
- 3) Client side of SMTP opens TCP connection with Bob’s mail server
- 4) SMTP client sends Alice’s message over the TCP connection
- 5) Bob’s mail server places the message in Bob’s mailbox
- 6) Bob invokes his user agent to read message



Sample SMTP interaction

S: 220 hamburger.edu
C: HELO crepes.fr
S: 250 Hello crepes.fr, pleased to meet you
C: MAIL FROM: <alice@crepes.fr>
S: 250 alice@crepes.fr... Sender ok
C: RCPT TO: <bob@hamburger.edu>
S: 250 bob@hamburger.edu ... Recipient ok
C: DATA
S: 354 Enter mail, end with "." on a line by itself
C: Do you like ketchup?
C: How about pickles?
C: .
S: 250 Message accepted for delivery
C: QUIT
S: 221 hamburger.edu closing connection

SMTP: final words

- ♦ SMTP uses persistent connections
- ♦ SMTP requires message (header & body) to be in 7-bit ASCII
- ♦ SMTP server uses CRLF . CRLF to determine end of message

Comparison with HTTP:


- ♦ HTTP: pull
- ♦ SMTP: push
- ♦ both have ASCII command/response interaction, status codes
- ♦ HTTP: each object encapsulated in its own response msg
- ♦ SMTP: multiple objects sent in multipart msg

Mail message format


SMTP: protocol for exchanging email messages

RFC 822: standard for text message format:

- ◆ header lines, e.g.,
 - » To:
 - » From:
 - » Subject:*different from SMTP commands!*
- ◆ body
 - » the “message”, ASCII characters only



Mail access protocols



- ♦ SMTP: delivery/storage to receiver's server
- ♦ Mail access protocol: retrieval from server
 - » POP3: Post Office Protocol [RFC 1939]
 - authorization (agent ↔ server) and download
 - » IMAP: Internet Mail Access Protocol [RFC 1730]
 - more features (more complex)
 - manipulation of stored messages on server
 - » HTTP: Gmail, Outlook, ProtonMail, etc.

POP3 protocol

authorization phase

- ♦ client commands:
 - » **user**: declare username
 - » **pass**: password
- ♦ server responses
 - » **+OK**
 - » **-ERR**

transaction phase, client:

- ♦ **list**: list message numbers
- ♦ **retr**: retrieve message by number
- ♦ **dele**: delete
- ♦ **quit**

```
S: +OK POP3 server ready
C: user bob
S: +OK
C: pass hungry
S: +OK User successfully logged on

C: list
S: 1 498
S: 2 912
S: .
C: retr 1
S: <message 1 contents>
S: .
C: dele 1
S: +OK Message 1 deleted.
C: retr 2
S: <message 1 contents>
S: .
C: dele 2
S: +OK Message 2 deleted.
C: quit
S: +OK POP3 server signing off
```

POP3 (more) and IMAP

More about POP3

- ◆ Previous example uses “download and delete” mode
- ◆ Bob cannot re-read e-mail if he changes client
- ◆ “Download-and-keep”: copies of messages on different clients
- ◆ POP3 is stateless across sessions

IMAP

- ◆ Keep all messages in one place: the server
- ◆ Allows user to organize messages in folders
- ◆ IMAP keeps user state across sessions:
 - » names of folders and mappings between message IDs and folder name

Domain Name System (DNS)

DNS: Domain Name System

People: many identifiers:

- » SSN, name, passport #

Internet hosts, routers:

- » IP address (32 bit) - used for addressing datagrams
- » “name”, e.g., ww.yahoo.com - used by humans

Q: map between IP addresses and name ?

Domain Name System:

- ♦ *distributed database* implemented in hierarchy of many *name servers*
- ♦ *application-layer protocol* host, routers, name servers to communicate to *resolve* names (address/name translation)
 - » note: core Internet function, implemented as application-layer protocol
 - » complexity at network’s “edge”

DNS

DNS services


- ◆ hostname to IP address translation
- ◆ host aliasing
 - » Canonical, alias names
- ◆ mail server aliasing
- ◆ load distribution
 - » replicated Web servers: set of IP addresses for one canonical name

Why not centralize DNS?

- ◆ single point of failure
- ◆ traffic volume
- ◆ distant centralized database
- ◆ maintenance

➔ would not *scale*!

Distributed, Hierarchical Database

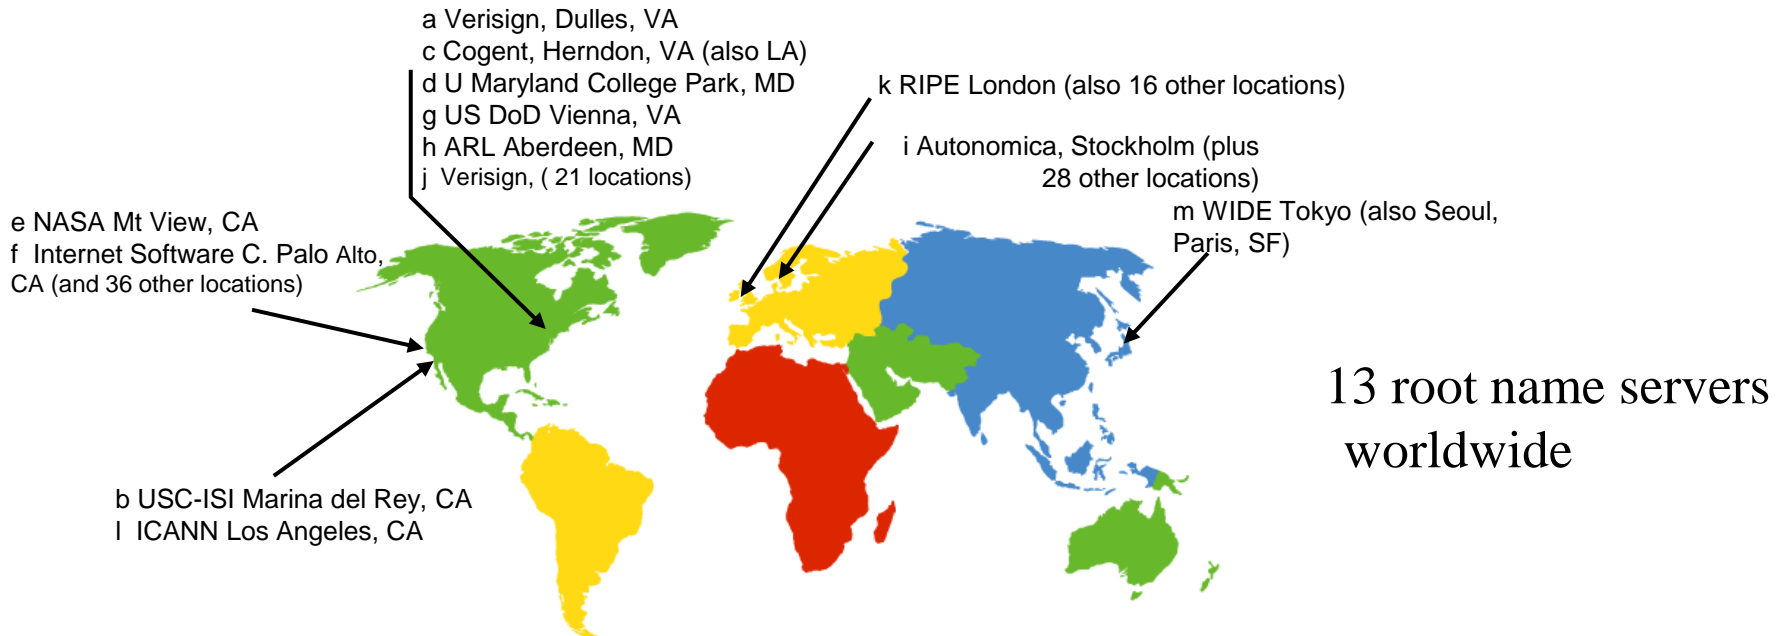


Client wants IP for www.amazon.com; 1st approx:

- ◆ client queries a [root server](#) to find [com DNS server](#)
- ◆ client queries [com DNS server](#) to [get amazon.com DNS server](#)
- ◆ client queries [amazon.com DNS server](#) to get [IP address for \[www.amazon.com\]\(http://www.amazon.com\)](#)

DNS: Root name servers

- ♦ contacted by local name server that can not resolve name
- ♦ return references to top-level domain servers



TLD and Authoritative Servers

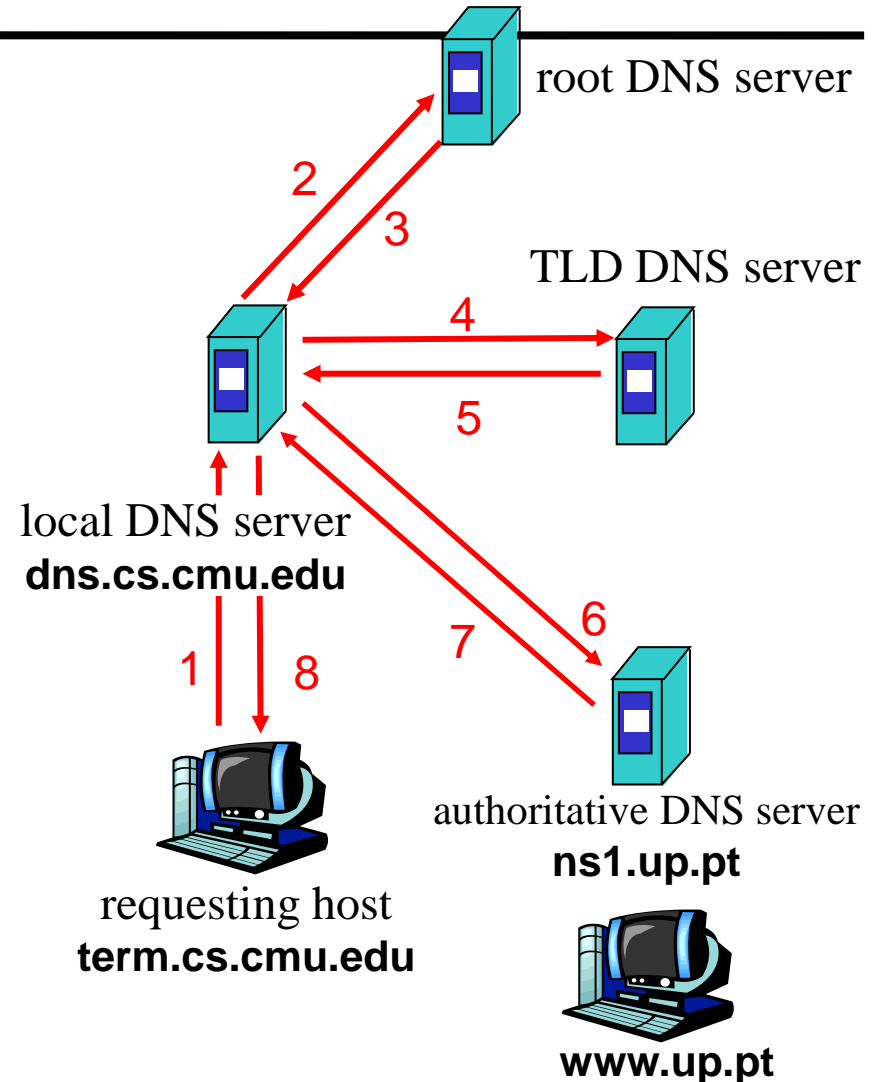
- ◆ **Top-level domain (TLD) servers:**
 - » responsible for com, org, net, edu, etc, and all top-level country domains uk, fr, ca, jp.
- ◆ **Authoritative DNS servers:**
 - » organization's DNS servers, providing authoritative hostname to IP mappings for organization's servers (e.g., Web, mail).
 - » can be maintained by organization or service provider

Local Name Server

- ◆ does not strictly belong to the hierarchy
- ◆ each ISP (residential ISP, company, university) has one
 - » also called “default name server”
- ◆ hosts make DNS queries to the local DNS server
 - » acts as proxy, querying the hierarchy to obtain the response
 - » caches results for efficiency

DNS name resolution example

- ◆ Host at `cs.cmu.edu` wants IP address for `www.up.pt`
- ◆ Makes **recursive query** to the local DNS server
 - » Response must be requested information or error
- ◆ The local DNS server resolves the name **iteratively**
 - » Makes **non-recursive queries** down the hierarchy
 - » Responses can be references to other DNS servers: *"I don't know, but ask this server..."*
- ◆ The local DNS server sends the final result to the client
 - » Caches final and intermediate results



DNS records

DNS: distributed database storing resource records (**RR**)

RR format: (**name**, **value**, **type**, **ttl**)

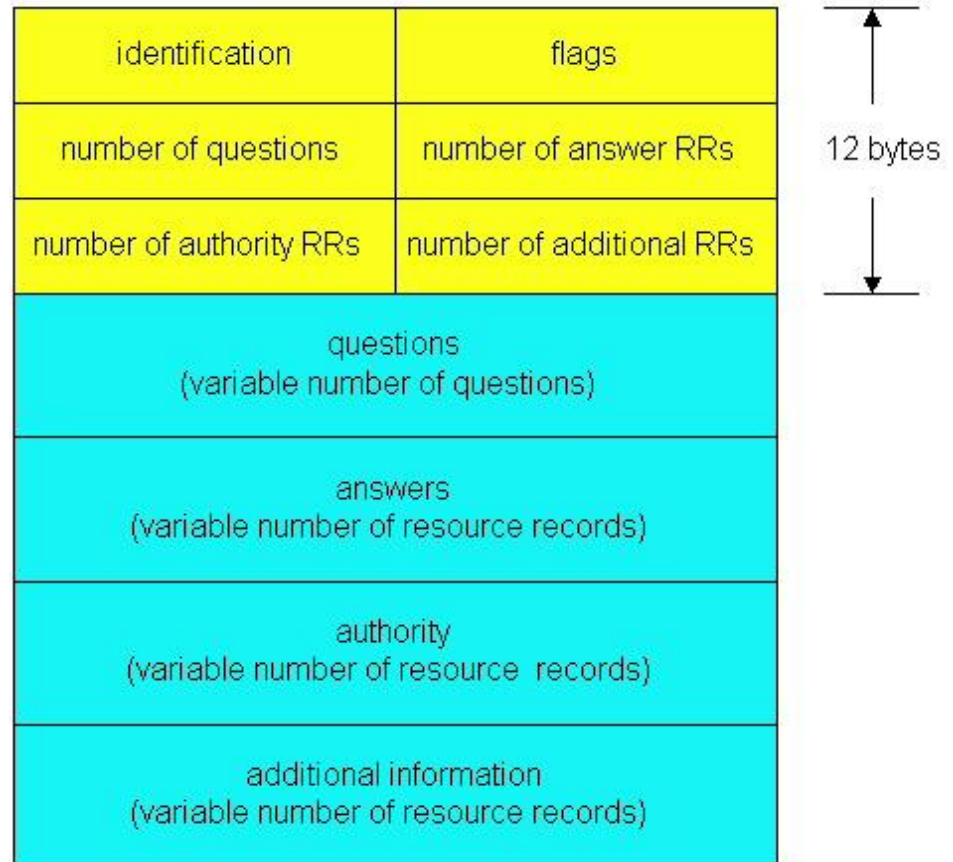
- ◆ Type=A
 - » **name** is hostname
 - » **value** is IP address
- ◆ Type=NS
 - » **name** is domain (e.g. foo.com)
 - » **value** is hostname of authoritative name server for this domain
- ◆ Type=CNAME
 - » **name** is alias name for some “canonical” (the real) name
www.ibm.com is really
servereast.backup2.ibm.com
 - » **value** is canonical name
- ◆ Type=MX
 - » **value** is name of mailserver associated with **name**

DNS protocol, messages

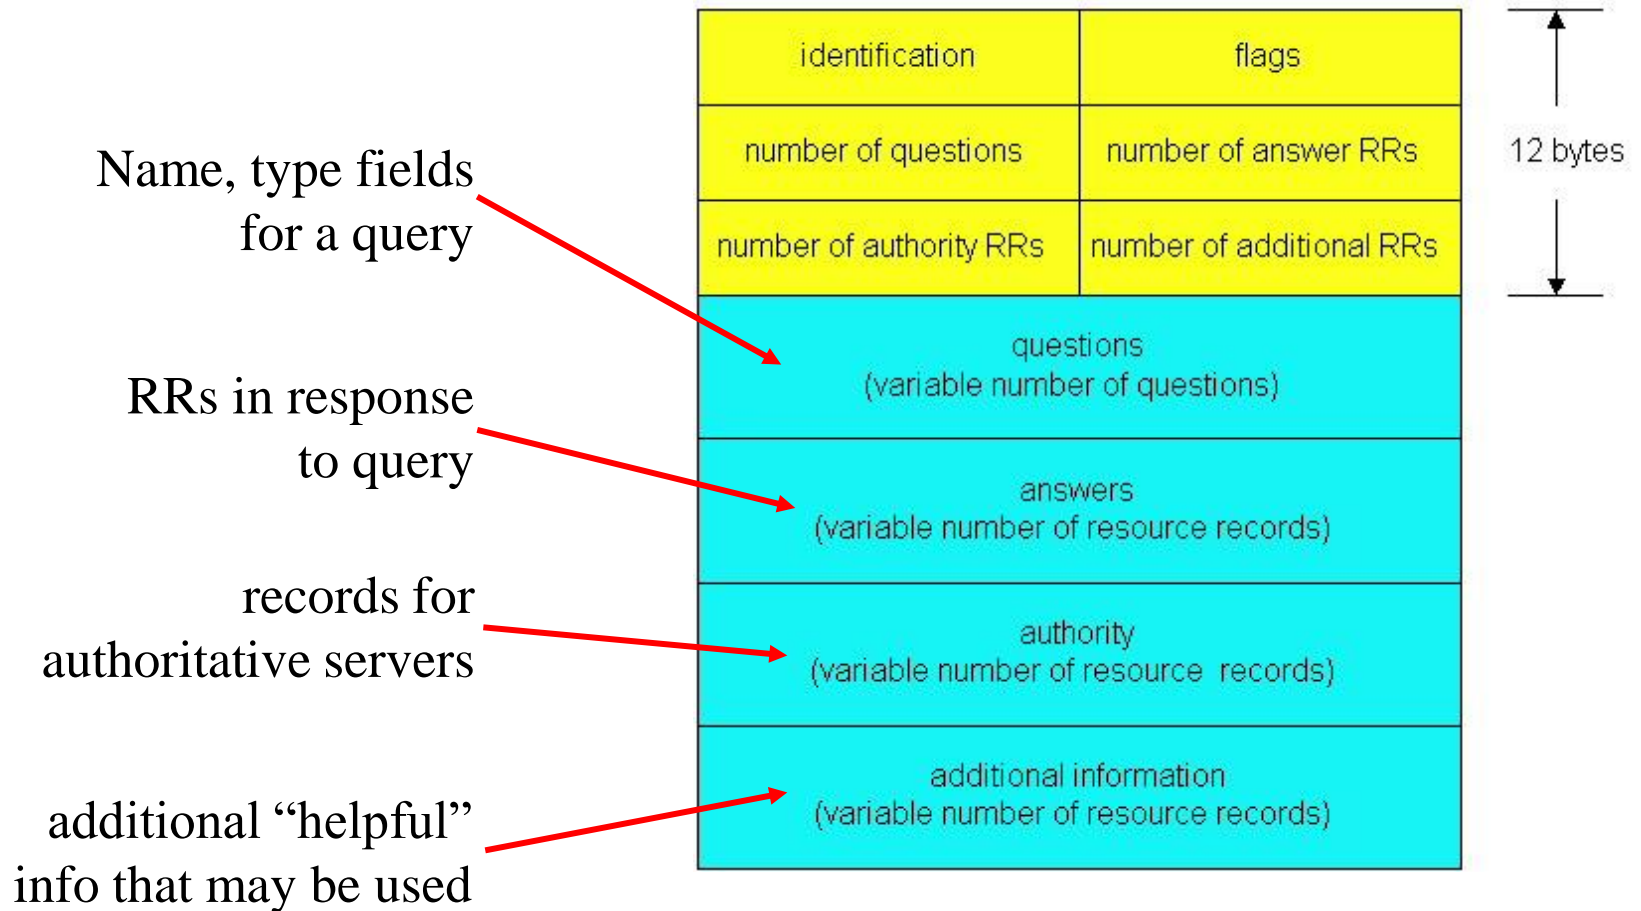
DNS protocol : *query* and *reply* messages, both with same *message format*

DNS message header:

- ♦ **identification**: 16 bit # for query, reply uses the same #
- ♦ **flags**:
 - » query or reply
 - » recursion desired
 - » recursion available
 - » reply is authoritative



DNS protocol, messages



Example of DNS response

```
[myself@localhost ~]$ dig mit.edu
```

```
; <<>> DiG 9.5.0-P1 <<>> mit.edu
```

```
;; global options: printcmd
```

```
;; Got answer:
```

```
;; ->>HEADER<<- opcode: QUERY, status: NOERROR, id: 28372
```

```
;; flags: qr rd ra; QUERY: 1, ANSWER: 1, AUTHORITY: 3, ADDITIONAL: 3
```

```
;; QUESTION SECTION:
```

```
mit.edu.                IN      A
```

```
;; ANSWER SECTION:
```

```
mit.edu.                60      IN      A      18.7.22.69
```

```
;; AUTHORITY SECTION:
```

```
mit.edu.                11809   IN      NS      STRAWB.mit.edu.
```

```
mit.edu.                11809   IN      NS      BITSY.mit.edu.
```

```
mit.edu.                11809   IN      NS      W20NS.mit.edu.
```

```
;; ADDITIONAL SECTION:
```

```
BITSY.mit.edu.          363     IN      A      18.72.0.3
```

```
W20NS.mit.edu.          3667    IN      A      18.70.0.160
```

```
STRAWB.mit.edu.         3667    IN      A      18.71.0.151
```


Inserting records into DNS

- ♦ example: new startup “Network Utopia”
- ♦ register name networkutopia.com at *DNS registrar* (e.g., Network Solutions)
 - » provide name(s) and IP address(es) of authoritative name server(s)
 - » registrar inserts two RRs into com TLD server:

```
networkutopia.com NS dns1.networkutopia.com  
dns1.networkutopia.com A 212.212.212.1
```

- ♦ in the authoritative server, create records for names in the domain
 - » A record for www.networkutopia.com
 - » MX record for networkutopia.com
 - » ...

Q: How do people get IP address of your Web site?